

Sift — Full-Codebase Evaluation Report

Scope: evaluate all production code, test and correct bugs, verify implementation against `MASTER_SPEC` and `ARCHITECTURE`, apply corrections/optimizations/improvements. **Date:** 2026-07-03 · **Branch:** `claude/creator-content-search-plan-h6o6kf`

1. Method

- Baseline verification** — full backend + frontend test suites, lint, both web builds, migration chain.
- Adversarial code review** (independent reviewer) of the one wave not previously reviewed (the deferred-items commit: cost metering, TikTok per-account capability, DM lock, injection sweep, pseudonym secret).
- Spec-conformance audit** (independent auditor) — every claim in `MASTER_SPEC` Part IV and every endpoint in `ARCHITECTURE` §3b checked against the actual code: all 67 `/agent` routes enumerated, every named function in all 18 agent modules verified non-trivial, ~12 load-bearing features semantically spot-checked, UI claims verified against the Qwik routes.
- Fixes, optimizations, improvements** applied in two committed, tested waves.

2. Baseline (before this evaluation)

311 backend + 9 frontend tests green · ruff clean · client + server builds green · migrations 0001–0004 applying. No pre-existing failures.

3. Findings

Adversarial review (deferred-items wave) — no HIGH; 4 MED/LOW, all fixed

Sev	Finding	Fix
MED	The anonymous-answer ceiling metered <code>op="answer"</code> from three surfaces (fan page, DM pipeline, dev API) — heavy legitimate DM/API volume would degrade the fan page with zero abuse	<code>generate_answer</code> takes an <code>op</code> tag; surfaces now meter as <code>answer</code> / <code>dm_answer</code> / <code>api_answer</code> / <code>eval_answer</code> ; the ceiling measures only the anonymous page
MED	Per-search ceiling check fetched + JSON-decoded every month-to-date event row in Python (linear growth with traffic)	Aggregation pushed into SQL (<code>SUM(payload->'tokens')</code>) + <code>op</code> filter in the WHERE, on the existing <code>(creator_id, kind, ts)</code> index
MED	<code>sync_native</code> could persist a <code>PlatformAccount</code> grant contradicting actual behavior (explicit full-loop connector, stored archive)	Grant + capabilities now persisted from the connector actually doing the sync
LOW/ MED	<code>get_connector</code> 's <code>try/except TypeError</code> could mask a genuine constructor bug as "no full_loop param" and silently downgrade a granted account	Replaced with <code>inspect.signature</code> parameter check

Verified-correct by the reviewer: metering exception flow, advisory-lock transaction span (no mid-dispatch commits), FakeLLM keyword matching unaffected by the injection guards, migration chain integrity.

Spec-conformance audit — docs substantially accurate; one systemic gap, closed

All 67 agent routes exist and match; zero stubs/TODOs in product code; the eval gate, DM compliance bounds, Studio UI claims, and test counts all verified exactly. The one systemic finding: **six features existed as implemented, tested library code with no production caller**. Per the goal ("implemented according to spec"), five were wired and one honestly re-marked:

Feature (spec §)	Now reachable via
Bounded clarifying question (§05, v1.5)	Public search response on no-strong-answer + fan-page either/or chips (zero-JS ? q= refinement)
No-answer waitlist / email-capture bridge (§05/§13, v1.5)	Fan-page email form → <code>POST /api/{handle}/waitlist</code> , cohort size shown back
Saved playbooks (§05, v1.5)	<code>match_playbook</code> applied in the comment → DM pipeline — matched template, still approval mode, skips an LLM call
Auto-affiliate injection (§06, v1.5)	Applied at draft approval (caption-time proxy — auto-posting stays deliberately avoided), with visible disclosure
Community delegation (§05, v2)	Candidate surfaced as a hint in the DM approval queue (the like itself stays manual — the official API has no comment-like endpoint)
Reply assistant (§05, v2)	Re-marked ▶ — it drives the agentic-DM thread path, which is itself honestly ▶

Also corrected in the docs: four rows marked "deferred" are actually **live opt-in scaffolds** (competitor radar, peer benchmarking, plagiarism scan — a real pgvector-similarity repost detector — and task export); three endpoint paths fixed; a fan-side surface catalog added (start-here / popular / topics / cohorts / feedback / save were undocumented); billing's location corrected (`cci_core.billing` , not providers). Both PDFs regenerated.

Own findings

- **Scheduler jobs had no retry policy** (previously flagged, never fixed) — a transient failure in a scheduled sync/radar/refresh silently dropped. Fixed: `DEFAULT_RETRY` on scheduled jobs; the DM dispatcher deliberately stays retry-free (next tick re-runs; the advisory lock makes overlap safe).
- **Digest delivery was a log stub** (documented gap). Implemented: `cci_core.mailer` (SMTP/STARTTLS behind config, log fallback — content never lost), `Creator.email` (migration 0005, settable via the admin API); digest + briefing now email when configured.

4. Changes shipped (2 commits)

Wave 1 — scheduler retries; digest email transport (+ `Creator.email` , migration 0005, 4 tests). **Wave 2** — all four review fixes; five spec wirings (+ fan-page waitlist form and clarify chips, `AnswerOut.clarify` , playbook path in `_prepare_dm` , affiliate injection in `decide_draft` , delegation hint in `dm-queue`); doc corrections; PDFs regenerated; 5 wiring tests + updated types/CSS.

5. Final verification

Check	Result
Backend tests	320 passed (was 311; +9 new)
Frontend tests (Vitest)	9 passed

Check	Result
Lint (ruff, all packages)	clean
Web builds (client + SSR server)	green
Migrations	0001 → 0005 apply cleanly
Spec conformance	Every v1 + v1.5 + v2 feature implemented and reachable ; v3 built or live-scaffolded except brand portal + marketplace (fully deferred, scale-gated); avoided-list still correctly unbuilt

6. Remaining known gaps (deliberate, documented)

- **Hosted OAuth consent screen** — tokens are operator-set; the flow downstream is built (environment-bound: needs live platform apps).
- **Agentic-DM thread dispatcher** (▶) — follow-up turns exist behind approval helpers; the dedicated dispatcher isn't wired.
- **Scale features** — brand portal, creator marketplace (fully deferred); competitor radar / peer benchmarking aggregates await a multi-creator panel.
- **Live-API quirks** — connectors are verified against faithful fakes; quota/auth edge cases unexercised offline.
- Pre-existing, accepted: Instagram sends happen inside the dispatch DB transaction (a crash mid-loop re-sends on the next run — concurrency dupes are locked out, crash dupes are not).

Verdict: the codebase is consistent with both specification documents, all discovered defects are fixed, and the docs now match the code in both directions.