

Sift — Agent Layer Roadmap

Companion to `IMPLEMENTATION_PLAN.md` (the v1 technical plan) and `ROADMAP_TO_FULL_APP.md` (productionization gaps). This document maps the **"Sift, as a Personal Agent" proposal (rev 2)** onto the code that exists today, so each feature reads as an increment, not a rewrite.

Sift v1 shows a creator what their audience wants. The agent layer acts on it — advise, execute, close — always behind the approval gate. **v1 ships unchanged**. The only thing the agent layer asks of v1 now is the cheap groundwork in §0.

Positioning shift: from "Search Console for your creator business" (passive, analytical) → **"Your AI chief of staff — Sift listens to your audience, tells you what they want, and helps you build it."**

Status legend: **built** built · **partial** partial (foundation exists) · — not started

The agent loop

The v1 loop has three audience-facing steps; the agent adds three that turn insight into action:

1 LISTEN	search + comments	← v1 built
2 SIFT	intent, filter, trends	← v1 built (+ §0 foundations)
3 SERVE	grounded answer to fan	← v1 built
4 ADVISE	actionable insight	← agent (v1.5)
5 EXECUTE	draft + links	← agent (v1.5–v2)
6 CLOSE	notify asker + measure	← agent (v2)

The loop closes when the original asker hears back — which seeds the next question.

§0 — Foundations (the ONLY thing that touches v1)

Build these into the v1 schema now. They cost almost nothing and are the difference between the agent layer being a feature set and being a rebuild.

Foundation	What it does	Release	Status in code today
Demand clustering & dedup	Merge semantically similar questions into one cluster ("142 people asked this"). Underpins every radar view.	v1	built <code>workers/radar.py</code> greedy centroid clustering
Intent / outcome event log	Log who asked what, what was served, what the creator did, and what happened next — the outcome edge.	v1	partial <code>events</code> table logs isolated events; not yet linked into question → served → action → result chains
Demand-item state machine	Each "useful" Radar card moves Idea → Drafting → Published → Loop-closed.	v1	partial <code>DemandTopic.creator_marked</code> (useful/not/made) only; no lifecycle states
Voice memory	Store approved replies/hooks and the creator's edits to drafts, so generated content stays on-voice.	v1 capture	— not started
Offer-awareness	Creator tags current products/offers/campaigns; Sift routes every CTA/affiliate to the right one.	v1.5	partial <code>products</code> table + approval exist; no "current offer/campaign" concept
Rules & preferences	Tone, taboo topics, escalation triggers, monetization priorities — the business rules the agent runs under.	v1.5	— not started

Concrete v1 work (small, additive): 1. Add a `status` lifecycle to `DemandTopic` (`idea` → `drafting` → `published` → `loop_closed`) alongside the existing `creator_marked`. 2. Extend the event log into linked **outcome records** (query/comment → answer served → creator action → measured result), so "did it work?" is answerable. 3. New **voice-memory** table capturing approved reply/hook text and draft edits.

Everything below waits for v1.5+.

§1 — Audience & Demand Intelligence ("What does my audience want?")

Demand Radar grows from a weekly list into a segmenting, emotion-reading chief of staff.

Feature	Release	Builds on
The Briefing — pushed daily/weekly digest with ready hooks, a re-promotion pick, open gaps, trend alerts, urgent replies, one fully drafted post	v1.5	partial <code>workers/digest.py</code> (currently a plain weekly digest)
Content gap map — answered well / partially / contradicted / never addressed	v1.5	partial radar coverage check already flags gaps
Persona segmentation — cluster audience by behaviour, per-segment feeds	v2	partial clustering infra
Sentiment & emotion — anxiety / confusion / excitement / frustration	v2	—
Trend detection — rising themes + resurfacing old posts; weekly default, spike alerts only above a volume threshold	v2	partial WoW deltas in radar
Cross-platform demand synthesis — newsletter BCC, forwarded DMs, podcast comments → "what's asked everywhere?"	v2	partial YouTube ingestion exists; multi-source intake new
Competitor demand radar — opt-in, aggregate, topic-level only	v3	—

§2 — Content Production ("What should I make next?")

Close the gap between a demand signal and a finished post — creator becomes editor, not writer.

Feature	Release	Builds on
Content briefs — one-pager from a demand cluster: phrasing, coverage gaps, hook + CTA, source links	v1.5	partial radar cards + retrieval
Draft generator (script + hooks) — reel script + hooks grounded in the creator's content and the audience's literal words	v1.5	partial LLM provider + retrieval
Creator recall search — creator-only "where did I say that?" with exact sentence + reuse ideas	v1.5	built retrieval layer (new surface, mostly reuse)
Multi-format repurposing + Sift & Shift (YouTube topic → IG carousel from transcript)	v2	partial transcripts + LLM
Series builder — planned arc from repeated demand	v2	partial radar
Performance prediction — score a draft vs the creator's top-10% patterns	v3	— needs performance history
Content calendar / planner — cadence-aware queue, posting windows, saturation flags	v3	—
Thumbnail & visual intelligence	v3	—

§3 — Engagement & Inbox ("What should I reply to, right now?")

The inbox becomes a triage layer + light CRM; comment-to-DM grows into an approval-mode assistant.

Feature	Release	Builds on
Intent-labelled queue — DMs/comments/story replies sorted: content request, purchase intent, support, collab lead	v1.5	partial <code>intent.py</code> detection + <code>dm_jobs</code>
Bounded clarifying question — one either/or to disambiguate ("pre- or post-workout?"), not open chat	v1.5	partial retrieval confidence
No-answer waitlist — "want a heads-up when Alex covers this?" captures email, feeds gap map	v1.5	partial no-answer state already exists
Saved playbooks — "when I get this kind of question, do this" templates	v1.5	—
Reply assistant — drafts public/DM replies in approved tone, grounded in citations	v2	partial DM drafting exists
Agentic DM (approval mode) — multi-turn within Meta's one-message/7-day limit; bulk approve; "always approve this type"	v2	partial <code>dm.py</code> dispatcher + window/cap enforcement
Community delegation — like a follower's good answer instead of spending the DM quota	v2	—
"Not now, but..." queue — defer a valid question, draft when the time comes, re-contact compliantly	v2	partial state machine (\$0)
Loop-Closer — notify the original asker when their request is published (opted-in channel / 7-day window)	v2	partial outcome log (\$0) + state machine

§4 — Monetization & Revenue

Every question is declared intent; turn it into income.

Feature	Release	Builds on
Auto-affiliate injection — append approved affiliate link + UTM at caption time	v1.5	built product→post mapping + tracked redirects
Offer-aware CTAs — attach the right current offer to answers/recommendations	v1.5	partial offer-awareness (\$0)
Affiliate optimisation — A/B link placement, anchor text, pairings	v2	partial event tracking
Sponsor matchmaker & pitch — demand-as-proof, match brand programs, draft pitch	v2	partial radar aggregates
Dynamic pricing & offer insights — objections, bundle signals, price-sensitivity by segment	v3	—
Revenue forecasting & goals	v3	—
Brand partnership portal	v3 · scale	—

§5 — Brand & Reputation

Feature	Release	Builds on
Crisis / sentiment-shift detection — negative spikes, controversy, web mentions + response framework	v2	—
Voice consistency scoring — score contractor/ghostwriter content vs the voice model	v2	partial voice memory (\$0)
Plagiarism / unauthorised-use monitoring	v3	—

§6 — Growth & Strategy

Feature	Release	Builds on
Content strategy advisor — expand vs double-down, segment growth, evergreen/trend balance	v2	partial accumulated Sift data
Peer benchmarking — opt-in / anonymised aggregate	v3	—
Education & skill planning — 1/3/5-year development plan from performance patterns	v3	—

§7 — Operations & Ecosystem

Feature	Release	Builds on
Task / project integration — Radar card → tracked task (Notion/Asana/Linear)	v3	partial state machine (\$0)
Customer CRM — unified profiles, lifecycle marketing	v3	—
Team & role-based access — owner/creator/VA/contractor + approval workflows + audit trail	v3	partial ties to the auth work in ROADMAP_TO_FULL_APP \$1
Creator marketplace	v3 · scale	—
Public API & dev ecosystem	v3 · scale	partial FastAPI foundation

Gated on scale, not the calendar: the brand portal, marketplace, and public API only matter at multi-creator scale — pursue only once the single-creator loop is proven and repeatable.

§8 — What to deliberately AVOID

Stated to keep the roadmap honest.

Excluded	Why
Open-ended chat	Hallucination/brand risk. The bounded answer card captures most of the value safely.
Auto-posting without approval	Too much trust too early. Approval mode stays the default through v3.
Real-time trend alerts (single creator)	Low signal-to-noise. Revisit only above a query-volume threshold.
Native video editing	Scope creep. Production stays at scripts/hooks/clip references, never an editor.
Cross-platform identity resolution	Privacy-radioactive under GDPR. Do topic-level demand synthesis, never individual identity matching.

§9 — Guardrails that carry over (unchanged from v1)

Already enforced in the codebase; none relax because the system became proactive:

- **Retrieval correctness is the gate** — citation correctness must pass on a held-out set before any automation goes live (`eval/ harness`, `--gate`).
- **Voice is a post-retrieval rewrite** — the agent rewrites a correct, cited answer into the creator's tone; it never generates novel claims. Correctness first, style second.
- **Approval mode by default** — `dm_jobs` start `pending_approval` ; bulk approval and "always approve this type" come later as trust is earned.
- **DM limits shape Close** — one creator-initiated message per comment within 7 days; re-contact only via opted-in channels or inside that window (`dm.py`).
- **Privacy stays aggregate and opt-in** — demand is pseudonymized and clustered (`privacy.py`); cross-platform/competitor features are opt-in; deletion path + policy before public launch.

§10 — Sequencing

Phase	What lands
v1 • unchanged (ship)	Search-first archive, Demand Radar, comment-to-DM. Groundwork only: state machine, intent/outcome event log, demand clustering, voice-memory capture.
v1.5 • assist (next)	The Briefing, content gap map, content briefs, draft + hook generator, creator recall search, bounded clarifying, no-answer waitlist, intent-labelled inbox, saved playbooks, auto-affiliate injection, offer-aware CTAs, offer/rules tagging.
v2 • agent (then)	Reply assistant, agentic DM, community delegation, Loop-Closer, "not now" queue, persona segmentation, sentiment, trend detection, cross-platform synthesis, multi-format repurposing + Sift & Shift, series builder, sponsor matchmaker + pitch, affiliate optimisation, performance-loop closure, strategy advisor, crisis detection, voice scoring.
v3 • scale (later)	Performance prediction, content calendar, thumbnail intelligence, dynamic pricing, revenue forecasting, customer CRM, task/PM integration, team/roles, peer benchmarking, education/skill planning, competitor radar, plagiarism monitoring, brand portal, marketplace, public API.

Recommended immediate action

Per the proposal's own emphasis, the single highest-leverage move is to land the **\$0 v1 groundwork** now, while the schema is young:

1. **Demand-item state machine** on `DemandTopic` (`idea` → `drafting` → `published` → `loop_closed`).
2. **Outcome-linked event log** — extend `events` so a question can be traced to what was served, what the creator did, and the measured result.
3. **Voice-memory capture** — a table for approved replies/hooks and draft edits.

These three are cheap, additive, don't change v1 behaviour, and are what make the entire v1.5–v3 agent layer an increment rather than a rebuild. The rest is correctly deferred and gated behind the same v1 validation guardrails.

Release tags are a sequencing hypothesis, not commitments — each feature should pass the v1 validation gates before it ships.